

Introduction to Mobile Robotics

Probabilistic Robotics

Wolfram Burgard, Michael Krawez

UTN

Probabilistic Robotics

Key idea:

Explicit representation of uncertainty

(using the calculus of probability theory)

- Perception = state estimation
- Action = utility optimization


Axioms of Probability Theory

$P(A)$ denotes probability that proposition A is true.

- $0 \leq P(A) \leq 1$
- $P(\text{True}) = 1$ $P(\text{False}) = 0$
- $P(A \vee B) = P(A) + P(B) - P(A \wedge B)$

A Closer Look at Axiom 3

$$P(A \vee B) = P(A) + P(B) - P(A \wedge B)$$



Using the Axioms

$$P(A \vee \neg A) = P(A) + P(\neg A) - P(A \wedge \neg A)$$

$$P(\textit{True}) = P(A) + P(\neg A) - P(\textit{False})$$

$$1 = P(A) + P(\neg A) - 0$$

$$P(\neg A) = 1 - P(A)$$

Discrete Random Variables


- X denotes a random variable
- X can take on a countable number of values in $\{x_1, x_2, \dots, x_n\}$
- $P(X=x_i)$ or $P(x_i)$ is the probability that the random variable X takes on value x_i
- $P(\cdot)$ is called probability mass function
- E.g., $P(Room) = \langle 0.7, 0.2, 0.08, 0.02 \rangle$

Continuous Random Variables

- X takes on values in the continuum.
- $p(X=x)$ or $p(x)$ is a **probability density function**

$$P(x \in [a, b]) = \int_a^b p(x) dx$$

- E.g.



“Probability Sums up to One”

Discrete case

$$\sum_x P(x) = 1$$

Continuous case

$$\int_X P(x) dx = 1$$

Joint and Conditional Probability

- $P(X=x \text{ and } Y=y) = P(x,y)$
- If X and Y are independent then
$$P(x,y) = P(x) P(y)$$
- $P(x / y)$ is the probability of x given y
$$P(x / y) = P(x,y) / P(y)$$
$$P(x,y) = P(x / y) P(y)$$
- If X and Y are independent then
$$P(x / y) = P(x)$$

Law of Total Probability

Discrete case

$$P(x) = \sum_y P(x | y)P(y)$$

Continuous case

$$p(x) = \int p(x | y)p(y)dy$$

Marginalization

Discrete case

$$P(x) = \sum_y P(x, y)$$

Continuous case

$$p(x) = \int p(x, y) dy$$

Bayes Formula

$$P(x, y) = P(x | y)P(y) = P(y | x)P(x)$$

$$P(x | y) = \frac{P(y | x)P(x)}{P(y)}$$

$$\frac{\text{Likelihood} * \text{Prior}}{\text{Evidence}}$$

Normalization

$$P(x | y) = \frac{P(y | x)P(x)}{P(y)}$$

- At the same time: $P(y) = \sum_x P(y | x)P(x)$

$$P(x | y) = \frac{P(y | x)P(x)}{\sum_x P(y | x)P(x)}$$

- $P(y)$ is independent of x and thus constant for all x

$$P(x | y) = \eta P(y | x)P(x)$$

Bayes Rule with Background Knowledge


$$P(x \mid y, a) = \frac{P(y \mid x, a)P(x \mid a)}{P(y \mid a)}$$

Conditional Independence

- $P(x, y | z) = P(x | z)P(y | z)$
- Equivalent to $P(x | z) = P(x | z, y)$ and $P(y | z) = P(y | z, x)$
- But this does not necessarily mean $P(x, y) = P(x)P(y)$
- Marginal independence does not mean independence

Simple Example of State Estimation

- Suppose a robot obtains measurement z
- What is $P(open \mid z)$?



Causal vs. Diagnostic Reasoning

- $P(open \mid z)$ is **diagnostic**
- $P(z \mid open)$ is **causal**
- In some situations, **causal** knowledge is easier to obtain
count frequencies!
- Bayes rule allows us to use causal knowledge:

$$P(open \mid z) = \frac{P(z \mid open)P(open)}{P(z)}$$

Example

- $P(z \mid open) = 0.6$ $P(z \mid \neg open) = 0.3$
- $P(open) = P(\neg open) = 0.5$
- $P(open \mid z) = \frac{P(z|open)P(open)}{P(z)} = \frac{0.6*0.5}{0.6*0.5+0.3*0.5} = \frac{2}{3}$
- z raises the probability that the door is open

Combining Evidence

- Suppose our robot obtains another observation z_2
- How can we integrate this new information?
- More generally, how can we estimate $P(x \mid z_1, \dots, z_n)$?

Recursive Bayesian Updating

$$P(x \mid z_1, \dots, z_n) = \frac{P(z_n \mid x, z_1, \dots, z_{n-1})P(x \mid z_1, \dots, z_{n-1})}{P(z_n \mid z_1, \dots, z_{n-1})}$$

Markov assumption:

z_n is independent of z_1, \dots, z_{n-1} given we know x

$$\begin{aligned} P(x \mid z_1, \dots, z_n) &= \frac{P(z_n \mid x)P(x \mid z_1, \dots, z_{n-1})}{P(z_n \mid z_1, \dots, z_{n-1})} \\ &= \alpha P(z_n \mid x)P(x \mid z_1, \dots, z_{n-1}) \\ &= \alpha P(x) \prod_{i=1 \dots n} P(z_i \mid x) \end{aligned}$$

Example: Second Measurement

- $P(z_2 \mid open) = 0.25$ $P(z_2 \mid \neg open) = 0.3$
- $P(open \mid z_1) = \frac{2}{3}$

$$\begin{aligned} P(open \mid z_2, z_1) &= \frac{P(z_2 \mid open)P(open \mid z_1)}{P(z_2 \mid open)P(open \mid z_1) + P(z_2 \mid \neg open)P(\neg open \mid z_1)} \\ &= \frac{\frac{1}{4} \cdot \frac{2}{3}}{\frac{1}{4} \cdot \frac{2}{3} + \frac{3}{10} \cdot \frac{1}{3}} = \frac{\frac{1}{6}}{\frac{1}{6} + \frac{1}{10}} = \frac{\frac{1}{6}}{\frac{4}{15}} = \frac{5}{8} = 0.625 \end{aligned}$$

- z_2 lowers the probability that the door is open

Actions

- Often the world is **dynamic** since
 - **actions carried out by the robot,**
 - **actions carried out by other agents,**
 - or just the **time** passing bychange the world
- How can we **incorporate** such **actions**?

Typical Actions

- The robot **turns its wheels** to move
- The robot **uses its manipulator** to grasp an object
- Plants grow over **time** ...

- Actions are **never carried out with absolute certainty**
- In contrast to measurements, **actions generally increase the uncertainty**


Modeling Actions

- To incorporate the outcome of an action u into the current “belief”, we use the conditional pdf

$$P(x / u, x')$$


- This term specifies the pdf that **executing u changes the state from x' to x .**

Example: Closing the door



State Transitions

$P(x \mid u, x')$ for $u = \text{“close door”}$:



If the door is open, the action “close door” succeeds in 90% of all cases

Integrating the Outcome of Actions

Continuous case:

$$P(x | u) = \int P(x | u, x') P(x' | \text{X}) dx'$$

Discrete case:

$$P(x | u) = \sum P(x | u, x') P(x' | \text{X})$$

We will make an independence assumption to get rid of the u in the second factor in the sum.

Example: The Resulting Belief

$$\begin{aligned}P(\textit{closed} \mid u) &= \sum P(\textit{closed} \mid u, x')P(x') \\&= P(\textit{closed} \mid u, \textit{open})P(\textit{open}) + P(\textit{closed} \mid u, \textit{closed})P(\textit{closed}) \\&= \frac{9}{10} \cdot \frac{5}{8} + \frac{1}{1} \cdot \frac{3}{8} = \frac{15}{16}\end{aligned}$$

$$\begin{aligned}P(\textit{open} \mid u) &= \sum P(\textit{open} \mid u, x')P(x') \\&= P(\textit{open} \mid u, \textit{open})P(\textit{open}) + P(\textit{open} \mid u, \textit{closed})P(\textit{closed}) \\&= \frac{1}{10} \cdot \frac{5}{8} + \frac{0}{1} \cdot \frac{3}{8} = \frac{1}{16} \\&= 1 - P(\textit{closed} \mid u)\end{aligned}$$

Bayes Filters: Framework

- **Given:**

- Stream of observations z and action data u :

$$d_t = \{u_1, z_1, \dots, u_t, z_t\}$$


- **Sensor model** $P(z \mid x)$
- **Action model** $P(x' \mid u, x)$
- **Prior** probability of the system state $P(x)$

- **Wanted:**

- Estimate of the state X of a **dynamical system**
- The posterior of the state is also called **Belief**:

$$Bel(x_t) = P(x_t \mid u_1, z_1, \dots, u_t, z_t)$$

Markov Assumption



$$P(z_t \mid x_{0:t}, z_{1:t-1}, u_{1:t}) = P(z_t \mid x_t)$$
$$P(x_t \mid x_{1:t-1}, z_{1:t-1}, u_{1:t}) = P(x_t \mid x_{t-1}, u_t)$$

Underlying Assumptions

- Static world
- Independent noise
- Perfect model, no approximation errors

Bayes Filters

z = observation
 u = action
 x = state

$$\boxed{Bel(x_t)} = P(x_t \mid u_1, z_1, \dots, u_t, z_t)$$

Bayes $= \eta P(z_t \mid x_t, u_1, z_1, \dots, u_t) P(x_t \mid u_1, z_1, \dots, u_t)$

Markov $= \eta P(z_t \mid x_t) P(x_t \mid u_1, z_1, \dots, u_t)$

Total prob. $= \eta P(z_t \mid x_t) \int P(x_t \mid u_1, z_1, \dots, u_t, x_{t-1})$
 $P(x_{t-1} \mid u_1, z_1, \dots, u_t) dx_{t-1}$

Markov $= \eta P(z_t \mid x_t) \int P(x_t \mid u_t, x_{t-1}) P(x_{t-1} \mid u_1, z_1, \dots, u_t) dx_{t-1}$

Markov $= \eta P(z_t \mid x_t) \int P(x_t \mid u_t, x_{t-1}) P(x_{t-1} \mid u_1, z_1, \dots, z_{t-1}) dx_{t-1}$

$$\boxed{= \eta P(z_t \mid x_t) \int P(x_t \mid u_t, x_{t-1}) Bel(x_{t-1}) dx_{t-1}}$$

$$Bel(x_t) = \eta P(z_t | x_t) \int P(x_t | u_t, x_{t-1}) Bel(x_{t-1}) dx_{t-1}$$


1. Algorithm **Bayes_filter**($Bel(x), d$):
2. $\eta = 0$
3. If d is a **perceptual** data item z then
 4. For all x do
 5. $Bel'(x) = P(z | x) Bel(x)$
 6. $h = h + Bel'(x)$
 7. For all x do
 8. $Bel'(x) = h^{-1} Bel'(x)$
9. Else if d is an **action** data item u then
 10. For all x do
 11. $Bel'(x) = \int P(x | u, x') Bel(x') dx'$
12. Return $Bel'(x)$

Bayes Filters are Familiar!

$$Bel(x_t) = \eta P(z_t | x_t) \int P(x_t | u_t, x_{t-1}) Bel(x_{t-1}) dx_{t-1}$$





- Kalman filters
- Particle filters
- Hidden Markov models
- Dynamic Bayesian networks
- Partially Observable Markov Decision Processes (POMDPs)

Probabilistic Localization



Probabilistic Localization

$$Bel(x \mid z, u) = \alpha p(z \mid x) \int_{x'} p(x \mid u, x') Bel(x') dx'$$



Summary

- Bayes rule allows us to compute probabilities that are hard to assess otherwise.
- Under the Markov assumption, recursive Bayesian updating can be used to efficiently combine evidence.
- Bayes filters are a probabilistic tool for estimating the state of dynamic systems.